



a place of mind

THE UNIVERSITY OF BRITISH COLUMBIA

CPSC 259

Dynamic memory

No iClickers in this deck!

Dynamic memory

- Arrays declared as local variables must have a known size at compile time
 - but sometimes we don't know how much space we need until runtime
- Suppose we expect users to only need to store up to 1000 values, so we hard-code "1000" as an array size
 - What if user needs more?

Change code and recompile
 - What if user only needs 5?

Wastes memory
- If the value 1000 is hard-coded, this is hard to find and change
 - especially if used in multiple locations
- If hardcoded as a symbolic constant, still cannot change without recompiling

Memory management in C

- We have already seen how locally-declared variables are placed on the function call stack
 - allocation and release are managed automatically
- The available stack space is extremely limited
 - placing many large variables or data structures on the stack can lead to stack overflow
- Stack variables only exist as long as the function that declared them is running

```
void MyFunction() {  
    int i; // local variable created on stack  
} // local variable goes out of scope and is deleted
```

Dynamic memory allocation

- At run-time, we can request extra space on-the-fly, from the *memory heap*
- Request memory from the heap – "allocation"
- Return allocated memory to the heap (when we no longer need it) – "deallocation"
- Unlike stack memory, items allocated on the heap must be **explicitly freed** by the programmer

Dynamic memory allocation

malloc returns NULL if unsuccessful

- Function **malloc** returns a pointer to a memory block of at least size bytes:

```
ptr = (cast-type*) malloc(byte-size);
```

local variable

*malloc returns void**

use sizeof (....)





- Function **free** returns the memory block (previously allocated with **malloc**) and pointed to by **ptr** to the memory heap:

```
free(ptr);
```

- The system knows how many bytes need to be freed, provided we supply the correct address **ptr**

Heap example

```
int main() {  
    int a;  
    int *p = (int*) malloc(sizeof(int));  
    *p = 10;  
}
```



- If there is no free memory left on the heap, `malloc` will return a null pointer
- Note: `malloc` only allocates the space but does not initialize the contents.
 - Use `calloc` to allocate and clear the space to binary zeros

Allocating dynamic arrays


- Suppose we want to allocate space for exactly 10 integers in an array

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int* i;
    i = (int*) malloc(10*sizeof(int));
    if (i == NULL) {
        printf("Error: can't get memory...\n");
        exit(1); // terminate processing
    }

    i[0] = 3; // equivalent: *(i+0) = 3;
    i[1] = 16; // *(i+1) = 16;
    printf("%d", *i);
    ...
}
```

can be variable



Allocating dynamic arrays

From user input, variable array size

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int employees, index;
    double* wages;
    printf("Number of employees? ");
    scanf("%d", &employees);

    wages = (double*) malloc(employees * sizeof(double))
    if (!wages) { // equivalent: if (wages == NULL)
        printf("Error: can't get memory...\n");
    }

    printf("Everything is OK\n");
    ...
}
```

See `dma_examples.c`

Dangling pointers

- When we are done with an allocated object, we free it so that the system can reclaim (and later reuse) the memory

```
int main() {  
    int* i = (int*) malloc(sizeof(int));  
    *i = 5;  
    free(i);  
  
    printf("%d", *i);  
}
```

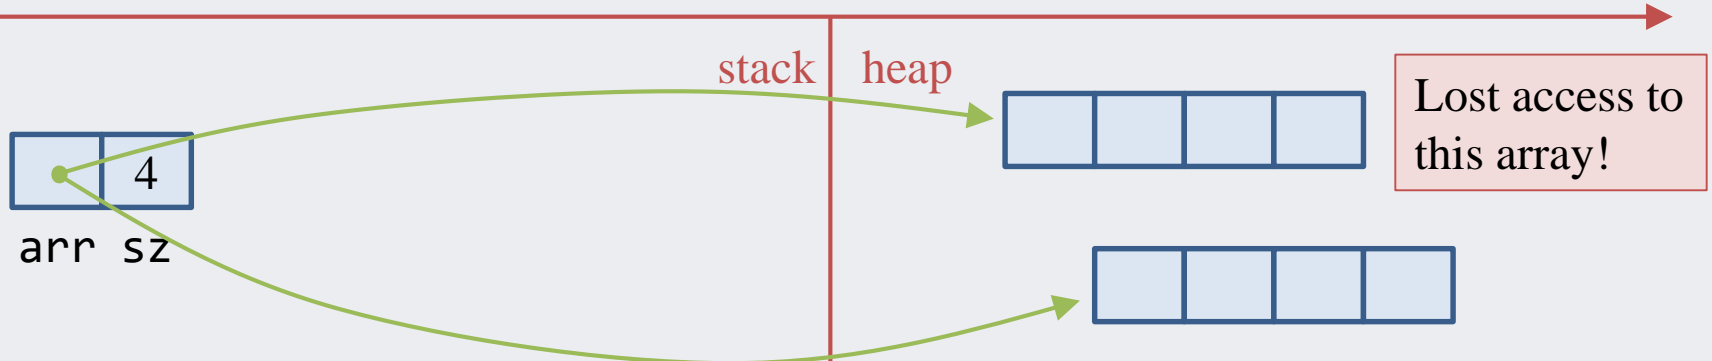
The space is marked as free, but the value remains until it is overwritten

- If the pointer continues to refer to the deallocated memory, it will behave unpredictably when dereferenced (and the memory is reallocated) – a **dangling pointer**
 - Leads to bugs that can be subtle and brutally difficult to find
 - So, set the pointer to **NULL** after freeing `i = NULL;`

Memory leaks

- If you lose access to allocated space (e.g. by reassigning a pointer), that space can no longer be referenced, or freed
 - And remains marked as allocated for the lifetime of the program

```
int* arr;  
int sz = 4;  
arr = (int*) malloc(sz*sizeof(int));  
arr[2] = 5;  
arr = (int*) malloc(sz*sizeof(int));  
arr[2] = 7;
```



See `memory_leak.c`

Exercise

- What is printed to the screen?
 - Also clearly identify memory leaks and dangling pointers

```
int w;  
int z;  
int* t = (int*) malloc(sizeof(int));  
int* y = (int*) malloc(sizeof(int));  
int* x = (int*) malloc(sizeof(int));  
  
*x = 3;  
*y = 5;  
z = *x + *y;  
w = *y;  
*x = z;  
free(x);  
*t = 2;  
y = &z;  
x = y;  
free(t);  
printf("*x=%d, *y=%d, z=%d, w=%d\n", *x, *y, z, w);
```

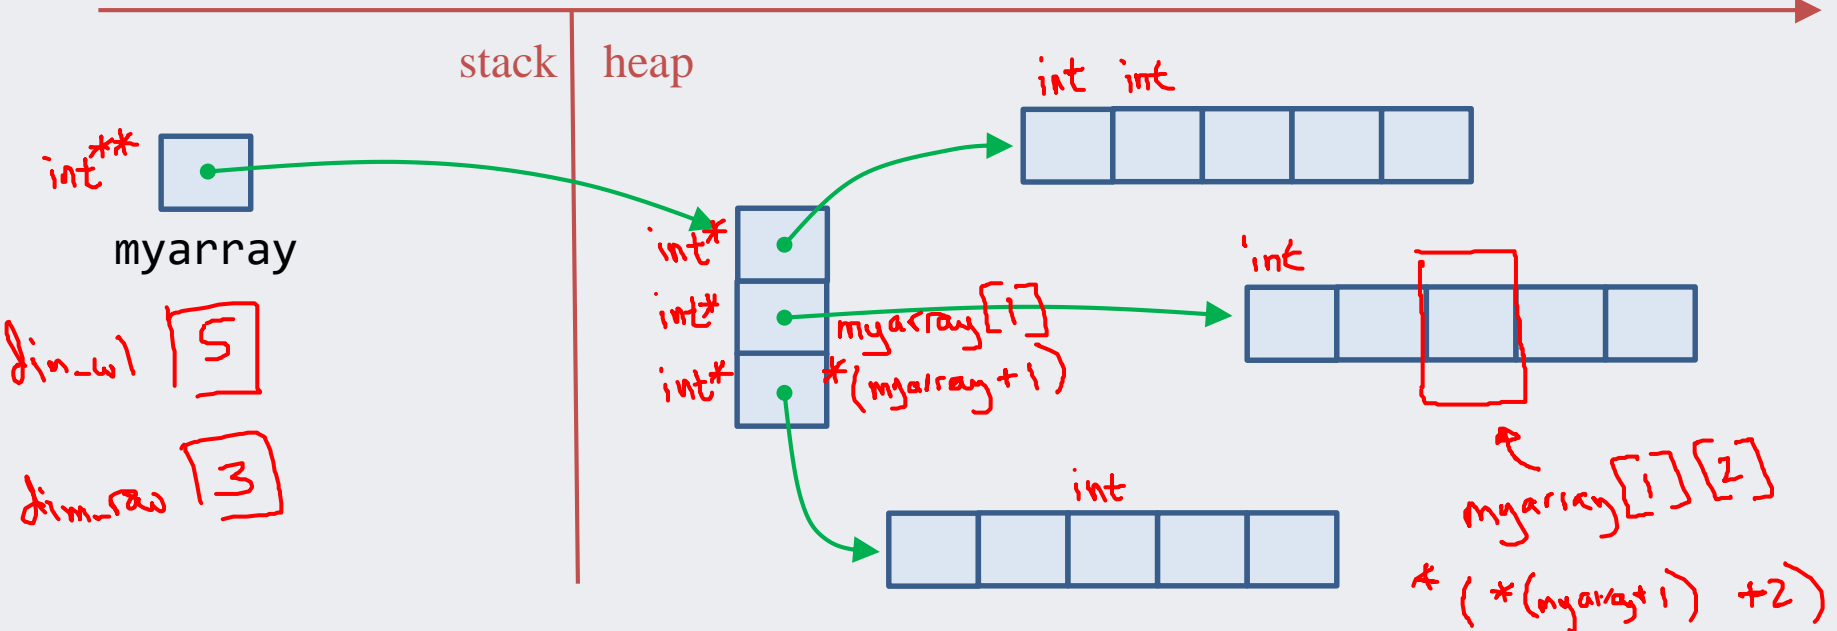
Dynamic allocation of a 2D array

See dma_2d.c for details

```
int dim_row = 3;  
int dim_col = 5;  
int** myarray; // pointer to a pointer
```

```
myarray = (int**) malloc(dim_row *  
                          sizeof(int*));  
for (int i=0 ; i< dim_row; i++) {  
    myarray[i] = (int*) malloc  
                  (dim_col * sizeof(int));  
}
```

(deallocate in reverse order)



Stack memory vs heap memory

- Stack
 - fast access
 - allocation/deallocation and space automatically managed
 - memory will not become fragmented
 - local variables only
 - limit on stack size (OS-dependent)
 - variables cannot be resized
- Heap
 - variables accessible outside declaration scope
 - no (practical) limit on memory size
 - variables can be resized
 - (relatively) slower access
 - no guaranteed efficient use of space
 - memory management is programmer's responsibility

Readings for this lesson

- Thareja
 - Appendices A, B, E
- Next class:
 - Thareja, Chapters 4 – 5